Gaussian LDA for Topic Models with **Word Embeddings**

Rajarshi Das, Manzil Zaheer, Chris Dyer • ACL • 2015

Continuous space **word embeddings** learned from large, unstructured corpora have been shown to be effective at capturing semantic regularities in language. In this paper we replace LDA's... (More)

20 30 View Paper Cite Save

Word Representations: A Simple and General Method for Semi-Supervised Learning

Joseph P. Turian, Lev-Arie Ratinov, Yoshua Bengio • ACL • 2010

If we take an existing supervised NLP system, a simple and general way to improve accuracy is to use unsupervised **word** representations as extra **word** features. We evaluate Brown clusters, Collobert... (More)

140 207 View Paper View Slides Cite Save

Dynamic **Word Embeddings**

Robert Bamler, Stephan Mandt • ICML • 2017

We present a probabilistic language model for time-stamped text data which tracks the semantic evolution of individual words over time. The model represents words and contexts by latent trajectories... (More)

4 14 View on ArXiv Cite Save

Morphological **Word-Embeddings**

Ryan Cotterell, Hinrich Schütze • HLT-NAACL • 2015

Linguistic similarity is multi-faceted. For instance, two words may be similar with respect to semantics, syntax, or morphology inter alia. Continuous **word-embeddings** have been shown to capture most... (More)

13 View Paper Cite Save

Deep contextualized **word** representations

Matthew E. Peters, Mark Neumann, +4 authors Luke S. Zettlemoyer • NAACL-HLT • 2018

We introduce a new type of deep contextualized **word** representation that models both (1) complex characteristics of **word** use (e.g., syntax and semantics), and (2) how these uses vary across linguistic... (More)

Results by year

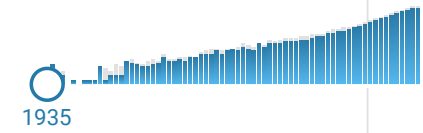

Slides related to word embeddings bias

Word Representations via Ga...
Luke Vilnis + 1 • Jul 18, 2018

Learning Word Embeddings f...
Yu-An Chung + 1 • Oct 17, 2018

Symmetric Pattern Based W...
Roy Schwartz + 2 • Apr 14, 201

Show More

Hiroaki Sakoe • 1978

This paper reports on an optimum dynamic programming (DP) based time-normalization algorithm for spoken **word** recognition. First, a general principle of time-normalization i s given using timewarping... (More)

395 225 | View Paper Cite Save

Gender **Bias** in Contextualized **Word Embeddings**

Jieyu Zhao, Tianlu Wang, Mark Yatskar, Ryan Cotterell, Vicente Ordonez, Kai-Wei Chang • 2019

In this paper, we quantify, analyze and mitigate gender **bias** exhibited in ELMo's contextualized **word** vectors. First, we conduct several intrinsic analyses and find that (1) training data for ELMo... (More)

View on ArXiv Cite Save

Word Segmentation : The Role of Distributional Cues

Jenny R. Saffran, Elissa L. Newport, Richard N. Aslin • 1996

One of the infant's first tasks in language acquisition is to discover the words **embedded** in a mostly continuous speech stream. This learning problem might be solved by using distributional cues to... (More)

54 29 | View Paper Cite Save

Reducing gender **bias** in **word embeddings**

Tuhin Chakraborty, Gabrielle Badie, Brett Rudder • 2016

Word embedding is a popular framework that represents text data as vectors of real numbers. These vectors capture semantics in language, and are used in a variety of natural language processing and... (More)

View Paper Cite Save

Learning **word embeddings** efficiently with noise-contrastive estimation

Andriy Mnih, Koray Kavukcuoglu • NIPS • 2013

Overview • Learning flexible **word** representations is the first step towards learning semantics. •The best current approach to learning **word embeddings** involves training a neural language model to... (More)

19 54 | View Paper Cite Save